

4-by-4 factorial design

- predictors as factors
- analysis of variance
- testing main effects
- group means and treatment differences
 - using emmeans
- other functions and packages for comparisons
 - aov
 - general contrasts using package `multcomp`
 - multiple comparisons using package `DescTools`
- using specific contrasts as coefficients

In an experiment with soybean, micronutrients were added to a fertilizer: copper (Cu) and/or manganese (Mn) (values in % of fertilizer). Yield was then measured (kg/acre).

```
soy = read.table("soybean.txt", header=T)
```

predictors as factors

Sometimes we will need to use `cu` as a numerical variable (to plot the data for instance), and sometimes as a factor with 4 categories (for the analysis), and similarly with `mn`. Below we define 2 new columns with the same data as in `cu` and `mn`, but considered as factors.

```
soy$cu_factor = factor(soy$cu)
soy$mn_factor = factor(soy$mn)
str(soy)
```

```
'data.frame': 32 obs. of 5 variables:
 $ cu      : int  1 1 1 1 3 3 3 3 5 5 ...
 $ mn      : int  20 50 80 110 20 50 80 110 20 50 ...
 $ yield   : int  1558 2003 2490 2830 1590 2020 2620 2860 1550 2010 ...
 $ cu_factor: Factor w/ 4 levels "1","3","5","7": 1 1 1 1 2 2 2 2 3 3 ...
 $ mn_factor: Factor w/ 4 levels "20","50","80",...: 1 2 3 4 1 2 3 4 1 2 ...
```


The experiments used 4 levels of cu addition, and 4 levels of mn addition, with 2 plots for each combination:

```
table(soy$cu, soy$mn)
```

| | 20 | 50 | 80 | 110 |
|---|----|----|----|-----|
| 1 | 2 | 2 | 2 | 2 |
| 3 | 2 | 2 | 2 | 2 |
| 5 | 2 | 2 | 2 | 2 |
| 7 | 2 | 2 | 2 | 2 |

Let's look at the data. We can also visualize and calculate the means at each combination:

```
layout(matrix(1:2,1,2))
plot(yield ~ cu, col=mn, pch=16, data=soy)
with(soy, interaction.plot(cu, mn, yield))
```




```
library(dplyr)
soy %>% group_by(cu, mn) %>% summarize(mean = mean(yield), sd = sd(yield))
```

```
# A tibble: 16 × 4
# Groups:   cu [4]
   cu    mn  mean    sd
  <int> <int> <dbl> <dbl>
1     1    20 1526.  44.5
2     1    50 2012.  12.7
3     1    80 2448.  59.4
4     1   110 2910. 113.
5     3    20 1606.  23.3
6     3    50 1979.  58.0
7     3    80 2644.  34.6
8     3   110 2897.  52.3
9     5    20 1579.  41.0
10    5    50 1994.  21.9
11    5    80 2566. 107.
12    5   110 2770.  27.6
13    7    20 1297.  43.8
14    7    50 1650. 156.
15    7    80 2336.  79.2
16    7   110 2628.   2.83
```

analysis of variance

Now to the formal analysis. It is extremely important to consider our two predictors as factors, to let each group have its own mean. Otherwise, and the mean for group cu=3 would be constrained to be half-way between the mean for group cu=1 and for group cu=5.

```
fit = lm(yield ~ cu_factor * mn_factor, data=soy)
plot(rstudent(fit)~fitted(fit), col=rgb(0,0,0, 0.5), pch=16,
     xlab="fitted values", ylab="ext. studentized residuals")
```



```
anova(fit)
```

Analysis of Variance Table

Response: yield

| | Df | Sum Sq | Mean Sq | F value | Pr(>F) |
|---------------------|----|---------|---------|----------|-----------|
| cu_factor | 3 | 443636 | 147879 | 32.2943 | 5.084e-07 |
| mn_factor | 3 | 8161105 | 2720368 | 594.0845 | < 2.2e-16 |
| cu_factor:mn_factor | 9 | 69190 | 7688 | 1.6789 | 0.1753 |
| Residuals | 16 | 73265 | 4579 | | |

In R, be really careful with this `anova` function: it returns the “type I” sums of squares, that is, when factors/predictors are added **sequentially**. Here, the design is balanced ($n=2$ in each treatment combination), so the type I and type III sums of squares are equal: the variance explained by each factor is the same whether it is added first or added last to the model. But in general, beware of the `anova` function: do not use it if you want to test a factor given that all others are in the model.

A safe way to get p-values based on type III sums of squares is using the `drop1` function, which drops only 1 term at a time:

```
drop1(fit, test="F")
```

Single term deletions

Model:

```
yield ~ cu_factor * mn_factor
```

| | Df | Sum of Sq | RSS | AIC | F value | Pr(>F) |
|---------------------|----|-----------|--------|--------|---------|--------|
| <none> | | | 73265 | 279.56 | | |
| cu_factor:mn_factor | 9 | 69190 | 142456 | 282.83 | 1.6789 | 0.1753 |

Unfortunately (or fortunately), `drop1` respects the hierarchy principle, and it won’t drop the term for “cu” if an interaction involving “cu” is present in the model. So here we only get the test for the interaction term.

testing main effects

Why can't we test for main effects easily in R? We didn't get p-values for main effects with `drop1`, and we don't get them with `anova` unless the design is balanced (type I, or sequential SS). Here is why:

- if there is an interaction:
 - all terms in this interaction should be considered as having an effect: no need for further testing.
 - the "main" effects are not meaningful: the main effect of a factor does not represent any of its simple effects
- if there is no interaction, then we can drop the interaction term from the model, and testing main effects becomes easy.

In our example the interaction is not significant, so we could consider dropping it to test each factor individually:

```
fit.noint = lm(yield ~ cu_factor + mn_factor, data=soy)
drop1(fit.noint, test="F")
```

Single term deletions

Model:

yield ~ cu_factor + mn_factor

| | Df | Sum of Sq | RSS | AIC | F value | Pr(>F) |
|-----------|----|-----------|---------|--------|---------|-----------|
| <none> | | | 142456 | 282.83 | | |
| cu_factor | 3 | 443636 | 586092 | 322.10 | 25.952 | 7.58e-08 |
| mn_factor | 3 | 8161105 | 8303561 | 406.93 | 477.406 | < 2.2e-16 |

However we changed the model: the interaction was removed, so the tests for the main effects have changed (notice the slightly different p-values): for instance, we tested "all $\alpha_i=0 \mid \mu, \beta_j$ " instead of "all $\alpha_i=0 \mid \mu, \beta_j, (\alpha\beta)_{ij}$ ". Many researchers would argue **against** removing the interaction from the model, even if it is not significant, because the experiment was designed as a 2-factorial CRD.

group means and treatment differences

The estimated coefficients tell us about the estimated group means, but getting all the group means is not straightforward.

```
summary(fit)$coefficients
```

| | Estimate | Std. Error | t value | Pr(> t) |
|-------------------------|----------|------------|------------|--------------|
| (Intercept) | 1526.5 | 47.84921 | 31.9023032 | 6.512077e-16 |
| cu_factor3 | 80.0 | 67.66900 | 1.1822252 | 2.543885e-01 |
| cu_factor5 | 52.5 | 67.66900 | 0.7758353 | 4.491629e-01 |
| cu_factor7 | -229.5 | 67.66900 | -3.3915086 | 3.726336e-03 |
| mn_factor50 | 485.5 | 67.66900 | 7.1746293 | 2.210244e-06 |
| mn_factor80 | 921.5 | 67.66900 | 13.6177567 | 3.226611e-10 |
| mn_factor110 | 1383.5 | 67.66900 | 20.4451074 | 6.813319e-13 |
| cu_factor3:mn_factor50 | -113.0 | 95.69842 | -1.1807928 | 2.549412e-01 |
| cu_factor5:mn_factor50 | -70.0 | 95.69842 | -0.7314645 | 4.750730e-01 |
| cu_factor7:mn_factor50 | -133.0 | 95.69842 | -1.3897826 | 1.836308e-01 |
| cu_factor3:mn_factor80 | 116.5 | 95.69842 | 1.2173660 | 2.411158e-01 |
| cu_factor5:mn_factor80 | 65.0 | 95.69842 | 0.6792171 | 5.067079e-01 |
| cu_factor7:mn_factor80 | 117.5 | 95.69842 | 1.2278155 | 2.372737e-01 |
| cu_factor3:mn_factor110 | -93.0 | 95.69842 | -0.9718029 | 3.456119e-01 |
| cu_factor5:mn_factor110 | -193.0 | 95.69842 | -2.0167522 | 6.082519e-02 |
| cu_factor7:mn_factor110 | -52.5 | 95.69842 | -0.5485984 | 5.908547e-01 |

- For example, the levels **cu=1** and **mn=20** do not appear in the coefficient names. They are the baseline levels, with values of 0 for all dummy variables that represent cu and mn. Do `model.matrix(fit)` to see these dummy variables. So the mean yield in the baseline group, cu=1 and mn=20 is the intercept: 1526.50. Its standard error is given as 47.85, which would allow us to get a confidence interval for instance (using a multiplier from the t-distribution with dfError = 16 df).
- For the group mean in group **cu=3 and mn=50**, it's more complicated: we need to add the intercept, the coefficient for cu=3, and coefficient for mn=50, and the interaction coefficient for "cu=3 and mn=50": $1526.50 + 80 + 485.5 + (-113) = 1979$. But the coefficient table does not give the SE for this estimate.

using emmeans

The package emmeans (<https://cran.r-project.org/web/packages/emmeans/>) is the best to do inference in treatment means, their pairwise differences, and other contrasts. It has great documentation, is very flexible (e.g. to models with random effects), and easier syntax than other packages to specify contrasts.

EMM stands for "estimated marginal means" (estimated from a fitted model), sometimes called adjusted means when the model has covariates. Let's look at treatment means first:

```
library(emmeans)
emmeans(fit, ~ mn_factor)
```

| mn_factor | emmean | SE | df | lower.CL | upper.CL |
|-----------|--------|------|----|----------|----------|
| 20 | 1502 | 23.9 | 16 | 1452 | 1553 |
| 50 | 1909 | 23.9 | 16 | 1858 | 1959 |
| 80 | 2498 | 23.9 | 16 | 2448 | 2549 |
| 110 | 2801 | 23.9 | 16 | 2750 | 2852 |

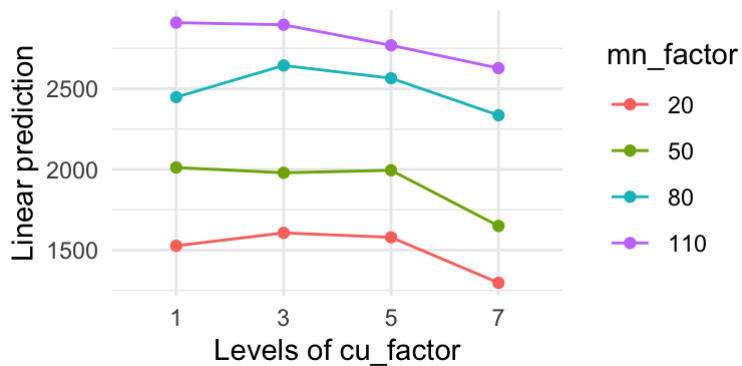
Results are averaged over the levels of: cu_factor
Confidence level used: 0.95

```
# emmeans(fit, ~ mn_factor | cu_factor) # same as below, presented differently
emmeans(fit, ~ mn_factor:cu_factor)
```


| mn_factor | cu_factor | emmean | SE | df | lower.CL | upper.CL |
|-----------|-----------|--------|------|----|----------|----------|
| 20 | 1 | 1526 | 47.8 | 16 | 1425 | 1628 |
| 50 | 1 | 2012 | 47.8 | 16 | 1911 | 2113 |
| 80 | 1 | 2448 | 47.8 | 16 | 2347 | 2549 |
| 110 | 1 | 2910 | 47.8 | 16 | 2809 | 3011 |
| 20 | 3 | 1606 | 47.8 | 16 | 1505 | 1708 |
| 50 | 3 | 1979 | 47.8 | 16 | 1878 | 2080 |
| 80 | 3 | 2644 | 47.8 | 16 | 2543 | 2746 |
| 110 | 3 | 2897 | 47.8 | 16 | 2796 | 2998 |
| 20 | 5 | 1579 | 47.8 | 16 | 1478 | 1680 |
| 50 | 5 | 1994 | 47.8 | 16 | 1893 | 2096 |
| 80 | 5 | 2566 | 47.8 | 16 | 2464 | 2667 |
| 110 | 5 | 2770 | 47.8 | 16 | 2668 | 2871 |
| 20 | 7 | 1297 | 47.8 | 16 | 1196 | 1398 |
| 50 | 7 | 1650 | 47.8 | 16 | 1548 | 1751 |
| 80 | 7 | 2336 | 47.8 | 16 | 2235 | 2437 |
| 110 | 7 | 2628 | 47.8 | 16 | 2527 | 2729 |

Confidence level used: 0.95


```
library(ggplot2)
emmip(fit, mn_factor ~ cu_factor) + theme_minimal() # ip = interaction plot
```



```
plot(emmeans(fit, ~ cu_factor))
```



```
plot(emmeans(fit, ~ cu_factor | mn_factor), by = "mn_factor")
```



by the way: why are all intervals of the same length?

Now, let's look at pairwise comparisons:

```
em_cu = emmeans(fit, ~ cu_factor)
pairs(em_cu) # default: Tukey
```

| contrast | estimate | SE | df | t.ratio | p.value |
|----------|----------|------|----|---------|---------|
| 1 - 3 | -57.6 | 33.8 | 16 | -1.703 | 0.3542 |
| 1 - 5 | -3.0 | 33.8 | 16 | -0.089 | 0.9997 |
| 1 - 7 | 246.5 | 33.8 | 16 | 7.285 | <.0001 |
| 3 - 5 | 54.6 | 33.8 | 16 | 1.614 | 0.3987 |
| 3 - 7 | 304.1 | 33.8 | 16 | 8.989 | <.0001 |
| 5 - 7 | 249.5 | 33.8 | 16 | 7.374 | <.0001 |

Results are averaged over the levels of: mn_factor

P value adjustment: tukey method for comparing a family of 4 estimates

```
pairs(em_cu, adjust = "none") # for LSD
```

| contrast | estimate | SE | df | t.ratio | p.value |
|----------|----------|------|----|---------|---------|
| 1 - 3 | -57.6 | 33.8 | 16 | -1.703 | 0.1079 |
| 1 - 5 | -3.0 | 33.8 | 16 | -0.089 | 0.9304 |
| 1 - 7 | 246.5 | 33.8 | 16 | 7.285 | <.0001 |
| 3 - 5 | 54.6 | 33.8 | 16 | 1.614 | 0.1260 |
| 3 - 7 | 304.1 | 33.8 | 16 | 8.989 | <.0001 |
| 5 - 7 | 249.5 | 33.8 | 16 | 7.374 | <.0001 |

Results are averaged over the levels of: mn_factor

```
pwpm(em_cu, adjust = "none") # pw = pairwise pm = p-value matrix
```

```

      1      3      5      7
1 [2224] 0.1079 0.9304 <.0001
3 -57.6 [2282] 0.1260 <.0001
5 -3.0   54.6 [2227] <.0001
7 246.5 304.1 249.5 [1978]
```

Row and column labels: cu_factor

Upper triangle: P values

Diagonal: [Estimates] (emmean)

Lower triangle: Comparisons (estimate) earlier vs. later

```
em_cuxmn = emmeans(fit, ~ cu_factor:mn_factor)[1:10] # first 10 only: to fit page width
pwpm(em_cuxmn)
```

```

      1 20      3 20      5 20      7 20      1 50      3 50      5 50      7 50      1 80      3 80
1 20 [1526] 0.9651 0.9980 0.0789 <.0001 0.0002 0.0001 0.7170 <.0001 <.0001
3 20 -80.0 [1606] 1.0000 0.0085 0.0006 0.0014 0.0009 0.9996 <.0001 <.0001
5 20 -52.5  27.5 [1579] 0.0185 0.0003 0.0007 0.0004 0.9842 <.0001 <.0001
7 20 229.5 309.5 282.0 [1297] <.0001 <.0001 <.0001 0.0025 <.0001 <.0001
1 50 -485.5 -405.5 -433.0 -715.0 [2012] 1.0000 1.0000 0.0019 0.0003 <.0001
3 50 -452.5 -372.5 -400.0 -682.0  33.0 [1979] 1.0000 0.0048 0.0001 <.0001
5 50 -468.0 -388.0 -415.5 -697.5  17.5 -15.5 [1994] 0.0031 0.0002 <.0001
7 50 -123.0 -43.0  -70.5 -352.5 362.5 329.5 345.0 [1649] <.0001 <.0001
1 80 -921.5 -841.5 -869.0 -1151.0 -436.0 -469.0 -453.5 -798.5 [2448] 0.1824
3 80 -1118.0 -1038.0 -1065.5 -1347.5 -632.5 -665.5 -650.0 -995.0 -196.5 [2644]
```

Row and column labels: cu_factor:mn_factor:.wgt.

Upper triangle: P values adjust = "tukey"

Diagonal: [Estimates] (emmean)

Lower triangle: Comparisons (estimate) earlier vs. later

Finally, let's consider a contrast other than a pairwise difference. Say, we want to see if the response at `cu=3` (averaged over the `mn` values) is the average between the mean at `cu=1` and at `cu=5` (again, averaged over the `mn` values). We find weak evidence against:

```
contrast(em_cu, list(cu3_vs1and5 = c(-1/2, 1, -1/2, 0)))
```

```

contrast      estimate      SE df t.ratio p.value
cu3_vs1and5      56.1 29.3 16   1.915  0.0735
```

Results are averaged over the levels of: mn_factor

other functions and packages for comparisons

Here are other tools that you may encounter, good to know about. But `emmeans` is the best so far, I think.

aov

The function `aov` gives us another way to fit the same model (and get the same ANOVA table), and has many tools for post-hoc group comparisons, like Tukey's honest significant differences. But it is limited to simple designs, and balanced designs.

```
fit.aov = aov(yield ~ cu_factor * mn_factor, data=soy)
anova(fit.aov)
```

Analysis of Variance Table

Response: yield

| | Df | Sum Sq | Mean Sq | F value | Pr(>F) |
|---------------------|----|---------|---------|----------|-----------|
| cu_factor | 3 | 443636 | 147879 | 32.2943 | 5.084e-07 |
| mn_factor | 3 | 8161105 | 2720368 | 594.0845 | < 2.2e-16 |
| cu_factor:mn_factor | 9 | 69190 | 7688 | 1.6789 | 0.1753 |
| Residuals | 16 | 73265 | 4579 | | |

```
TukeyHSD(fit.aov, "cu_factor")
```

Tukey multiple comparisons of means
95% family-wise confidence level


Fit: aov(formula = yield ~ cu_factor * mn_factor, data = soy)

\$cu_factor

| | diff | lwr | upr | p adj |
|-----|----------|------------|------------|-----------|
| 3-1 | 57.625 | -39.17618 | 154.42618 | 0.3541571 |
| 5-1 | 3.000 | -93.80118 | 99.80118 | 0.9997390 |
| 7-1 | -246.500 | -343.30118 | -149.69882 | 0.0000100 |
| 5-3 | -54.625 | -151.42618 | 42.17618 | 0.3986692 |
| 7-3 | -304.125 | -400.92618 | -207.32382 | 0.0000007 |
| 7-5 | -249.500 | -346.30118 | -152.69882 | 0.0000086 |

```
plot(TukeyHSD(fit.aov, "cu_factor"))
```

95% family-wise confidence level



Estimated group means are also easier to obtain after fitting the model with `aov` :

```
model.tables(fit.aov, type="means", se=TRUE)
```

Tables of means

Grand mean

2177.656

cu_factor

cu_factor

| | 1 | 3 | 5 | 7 |
|--|--------|--------|--------|--------|
| | 2224.1 | 2281.7 | 2227.1 | 1977.6 |

mn_factor

mn_factor

| | 20 | 50 | 80 | 110 |
|--|--------|--------|--------|--------|
| | 1502.2 | 1908.7 | 2498.5 | 2801.1 |

cu_factor:mn_factor

| | | mn_factor | | |
|-----------|--------|-----------|--------|--------|
| cu_factor | 20 | 50 | 80 | 110 |
| 1 | 1526.5 | 2012.0 | 2448.0 | 2910.0 |
| 3 | 1606.5 | 1979.0 | 2644.5 | 2897.0 |
| 5 | 1579.0 | 1994.5 | 2565.5 | 2769.5 |
| 7 | 1297.0 | 1649.5 | 2336.0 | 2628.0 |

Standard errors for differences of means

| | cu_factor | mn_factor | cu_factor:mn_factor |
|---------|-----------|-----------|---------------------|
| | 33.83 | 33.83 | 67.67 |
| replic. | 8 | 8 | 2 |

The last table gives SEs to calculate various things:

- the **LSD** (least significant difference) values: to compare the overall means among cu groups at the 5% level, for instance, the LSD is $33.83 * 2.12 = 71.72$. Group means that differ by more than 71.72 would be said statistically different (but recall that LSD is liberal). The 2.12 is the t multiplier for 95% confidence, from 16 df (dfError): `qt(.975, df=16) = 2.12`. So for instance, the means for `cu=1` and `cu=3` are not statistically different according to LSD: $2281.7 - 2224.1 = 57.6$, which is less than 71.72.
- test for **differences between means**: to compare the mean at `cu=3` versus `cu=1`, we would divide the observed difference in means, 57.6, by the standard error of this difference: 33.83 (last table in the output above). Note that this SE can be derived manually using the MSEError (4579) and the formula for the variance of a contrast: $\sqrt{4579 * (1/8 + 1/8)} = 33.83$. So our t-value is $57.6 / 33.83 = 1.702$, to be compared with a t-distribution with 16 df (error df): `pt(1.702, df=16, lower.tail=FALSE)*2 = 0.13`. Again, the difference is not significant.
- confidence interval for a **single mean**: we would follow a similar strategy. For the overall mean 2224.1 in groups `cu=1`, for example, the standard error of this mean is $\sqrt{4579 * (1/8)} = 23.92$. The t multiplier for 95% confidence is still 2.12 (16 df), so our confidence interval has limits $2224.1 \pm 2.12 * 23.92$.

- confidence interval for the group $cu=1$ and $mn=20$. Here we would use the estimated mean in that group: 1526.5. Its standard error is calculated on the basis of only 2 observation for that group:
 $\text{sqrt}(4579*(1/2)) = 47.85$.
- difference between the means in group “ $cu=1$ and $mn=20$ ” and group “ $cu=3$ and $mn=20$ ”: same strategy as before, but the standard error for the difference of interest is $\text{sqrt}(4579*(1/2 + 1/2))$, because each group only has 2 observations. In the output above, the last row of the last table is useful to determine the appropriate SE, based on which group(s) we are interested in and the sample size for that (those) groups.

We can also get the “effects” that we need to add to the overall mean (intercept in the model). Using the course notations, these effects are the α_i , β_j and $(\alpha\beta)_{ij}$ terms. Note that they do sum to zero, in each row and each column:

```
model.tables(fit.aov) # default is type="effects"
```

Tables of effects

```
cu_factor
cu_factor
  1      3      5      7
46.47 104.09  49.47 -200.03

mn_factor
mn_factor
 20      50      80     110
-675.4 -268.9  320.8  623.5

cu_factor:mn_factor
      mn_factor
cu_factor 20      50      80     110
  1 -22.22  56.78 -96.97  62.41
  3   0.16 -33.84  41.91  -8.22
  5  27.28  36.28  17.53 -81.09
  7  -5.22 -59.22  37.53  26.91
```

general contrasts using package `multcomp`

```
library(multcomp)
```

First, let’s look at a very specific contrast: at $mn=20$, is the response to $cu=3$ in between those at $cu=1$ and $cu=5$? Warning: the coefficient named `cu_factor3` is the difference between $cu=3$ and $cu=1$ (etc.), so the contrast coefficient is 0 for the intercept, below.

```
# coef(fit) # look at coefficients and what they mean, to define the contrast next
K = matrix(c(0,-2,1, rep(0,13)), 1)
t = glht(fit, linfct = K)
summary(t)
```

Simultaneous Tests for General Linear Hypotheses

```
Fit: lm(formula = yield ~ cu_factor * mn_factor, data = soy)
```

Linear Hypotheses:

| | Estimate | Std. Error | t value | Pr(> t) |
|--------|----------|------------|---------|----------|
| 1 == 0 | -107.5 | 117.2 | -0.917 | 0.373 |

(Adjusted p values reported -- single-step method)

testing main effect of cu using Tukey's HSD or Fisher's LSD:

```
lh_cu = glht(fit, linfct = mcp(cu_factor="Tukey", interaction_average=T))  
summary(lh_cu) # p-values: for Tukey's HSD
```

Simultaneous Tests for General Linear Hypotheses

Multiple Comparisons of Means: Tukey Contrasts

```
Fit: lm(formula = yield ~ cu_factor * mn_factor, data = soy)
```

Linear Hypotheses:

| | Estimate | Std. Error | t value | Pr(> t) |
|------------|----------|------------|---------|----------|
| 3 - 1 == 0 | 57.62 | 33.84 | 1.703 | 0.354 |
| 5 - 1 == 0 | 3.00 | 33.84 | 0.089 | 1.000 |
| 7 - 1 == 0 | -246.50 | 33.84 | -7.285 | <0.001 |
| 5 - 3 == 0 | -54.62 | 33.84 | -1.614 | 0.399 |
| 7 - 3 == 0 | -304.12 | 33.84 | -8.989 | <0.001 |
| 7 - 5 == 0 | -249.50 | 33.84 | -7.374 | <0.001 |

(Adjusted p values reported -- single-step method)

```
summary(lh_cu, test=univariate()) # no correction for multiple comparisons: Fisher LSD
```

Simultaneous Tests for General Linear Hypotheses

Multiple Comparisons of Means: Tukey Contrasts

```
Fit: lm(formula = yield ~ cu_factor * mn_factor, data = soy)
```

Linear Hypotheses:

| | Estimate | Std. Error | t value | Pr(> t) |
|------------|----------|------------|---------|----------|
| 3 - 1 == 0 | 57.62 | 33.84 | 1.703 | 0.108 |
| 5 - 1 == 0 | 3.00 | 33.84 | 0.089 | 0.930 |
| 7 - 1 == 0 | -246.50 | 33.84 | -7.285 | 1.83e-06 |
| 5 - 3 == 0 | -54.62 | 33.84 | -1.614 | 0.126 |
| 7 - 3 == 0 | -304.12 | 33.84 | -8.989 | 1.19e-07 |
| 7 - 5 == 0 | -249.50 | 33.84 | -7.374 | 1.57e-06 |

(Univariate p values reported)

applying Tukey's HSD to compare *all* 16 treatment means (try for yourself):

```
soy$cuxmn = with(soy, interaction(cu, mn, sep = "x"))
head(soy$cuxmn)
fit2 = lm(yield ~ cuxmn, data=soy)
lh2 = glht(fit2, linfct = mcp(cuxmn="Tukey"))
summary(lh2)
```

multiple comparisons using package DescTools

and aov fit (try for yourself):

```
library(DescTools)
PostHocTest(fit.aov, method = "lsd")
PostHocTest(fit.aov, method = "scheffe")
PostHocTest(fit.aov, method = "hsd")
```

using specific contrasts as coefficients

In our model `fit`, all levels except the first have their own coefficients. For each factor, the first level is the “base” level. The intercept corresponds to the mean yield when all the factors are at their “base” level. This parametrization of the design matrix and coefficient is said to use what's called, in R, the “treatment” contrasts. It facilitates pairwise comparisons between treatments.

```
options()$contrasts # default is "contr.treatment"
```

| unordered | ordered |
|-------------------|--------------|
| "contr.treatment" | "contr.poly" |

A different parametrization is with the “sum” contrasts, where the `cu` coefficients are α_i with $\sum \alpha_i = 0$ for example. Because of the constraint that α ’s (and β ’s, γ ’s etc) sum to 0, R does not estimate the last α_i , here the coefficient associated with level `cu=7` . It takes it as the negative sum of the others: $\alpha_7 = -(\alpha_1 + \alpha_3 + \alpha_5)$, using here the levels of `cu` to index the α coefficients.

```
fit.sum = lm(yield ~ cu_factor * mn_factor, data=soy,
             contrasts = list(cu_factor=contr.sum,
                             mn_factor=contr.sum))

summary(fit.sum)
```

```
Call:
lm(formula = yield ~ cu_factor * mn_factor, data = soy, contrasts = list(cu_factor = con
tr.sum,
    mn_factor = contr.sum))

Residuals:
    Min       1Q   Median       3Q      Max
-110.50  -31.12    0.00   31.12  110.50

Coefficients:
                Estimate Std. Error t value Pr(>|t|)
(Intercept)      2177.6562    11.9623  182.043 < 2e-16
cu_factor1         46.4687    20.7193   2.243 0.039426
cu_factor2        104.0938    20.7193   5.024 0.000125
cu_factor3         49.4688    20.7193   2.388 0.029643
mn_factor1        -675.4063    20.7193 -32.598 4.63e-16
mn_factor2       -268.9062    20.7193 -12.979 6.56e-10
mn_factor3        320.8438    20.7193  15.485 4.74e-11
cu_factor1:mn_factor1 -22.2188    35.8869  -0.619 0.544545
cu_factor2:mn_factor1   0.1563    35.8869   0.004 0.996580
cu_factor3:mn_factor1  27.2813    35.8869   0.760 0.458191
cu_factor1:mn_factor2  56.7813    35.8869   1.582 0.133162
cu_factor2:mn_factor2 -33.8438    35.8869  -0.943 0.359670
cu_factor3:mn_factor2  36.2812    35.8869   1.011 0.327067
cu_factor1:mn_factor3 -96.9687    35.8869  -2.702 0.015705
cu_factor2:mn_factor3  41.9062    35.8869   1.168 0.260023
cu_factor3:mn_factor3  17.5312    35.8869   0.489 0.631814

Residual standard error: 67.67 on 16 degrees of freedom
Multiple R-squared:  0.9916,    Adjusted R-squared:  0.9838
F-statistic: 126.3 on 15 and 16 DF,  p-value: 1.024e-13
```

With these “sum” contrasts, the intercept represents the mean of all group means, which is very nice. But getting individual group means is much more tedious. For example, to get the mean at `cu=3` (second level) and `mn=80` (third level):

```
2177.6562 + 104.0938 + 320.8438 + 41.9062 # 2644.5
```

Group means involving the last level for one or more factors are more complicated. For example, the mean yield at `cu=7` (last level) and `mn=20` (first level) is:

```
2177.6562 -(46.4687+104.0938+49.4688) + -675.4063 -(-22.2188+0.1563+27.2813) # 1297
```

An alternative way to fit the model with the “sum” contrasts is to change the global default contrast settings, and the “sum” contrasts will be use for all factors, now and later:

```
options(contrasts = c("contr.sum", "contr.poly"))  
fit.sum = lm(yield ~ cu_factor * mn_factor, data=soy)  
summary(fit.sum) # same as above
```