

# One-factor CRD with subsampling

## the data

An experiment was conducted as a completely randomized design with sub-sampling: there were 4 treatments, and 4 plots for each treatment. Within each plot 3 measurements (subsamples) were taken. In the data file, the column “core” indicates the subsample number. The column “y” contains the response.


```
dat = read.csv("crdsub.csv", colClasses=c("factor","factor","factor","numeric"))
head(dat, n=4)
```

	trt	plot	core	y
1	1	1	1	12.6
2	1	1	2	11.9
3	1	1	3	12.3
4	1	2	1	13.0

```
with(dat, table(trt,plot))
```

	plot			
trt	1	2	3	4
1	3	3	3	3
2	3	3	3	3
3	3	3	3	3
4	3	3	3	3

```
ggplot(dat, aes(y=y, x=trt, shape=plot, color=plot)) + geom_point() +theme_minimal()
```



# manual analysis with fixed-effects

For the analysis, we can start with a fixed-effect model to get the ANOVA table. The random effect for “plot” should be indicated with `plot`, because the plot values go from 1 to 4 only, and plot 1 means totally different plots across treatments (say). So we use `trt:plot` to get a different plot value for each plot (e.g. `1:1` for plot 1 in treatment 1, and `2:1` for plot 1 in treatment 2).

```
fit1 = lm(y ~ trt + trt:plot, dat)
anova(fit1) # warning: wrong F test for trt, because fixed-effect model
```

## Analysis of Variance Table

```
Response: y
      Df Sum Sq Mean Sq F value    Pr(>F)
trt      3  29.407   9.8024   76.507 1.097e-14
trt:plot 12  17.386   1.4488   11.308 2.342e-08
Residuals 32   4.100   0.1281
```

**Warning** about the output above: the F-test is wrong for the fixed effects (treatment), because the plot effects are considered fixed here instead of random. The test for plot variation is correct though.

Since this is a balanced design, so we can use the SS and MS from the `anova` function above (which returns type 1 SSs, but type 1 = type 3 SS with balanced designs). So here is a correct test for treatment differences:

```
f = 9.8024 / 1.4488; f          # f = 6.765875
pf(f, df1=3, df2=12, lower.tail=F) # p-value = 0.006365391
(1.4488 - 0.1281)/3          # sigma2_plot = 0.4402333
```

**Warning** again: the fixed effects model has wrong SEs for treatment means (e.g. intercept = mean of treatment 1 in plot 1 here) and wrong SEs for treatment differences.

```
head(summary(fit1)$coefficients, n=14)
```

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	12.2666667	0.2066599	59.3567689	2.608221e-34
trt2	0.1000000	0.2922613	0.3421596	7.344680e-01
trt3	-0.3666667	0.2922613	-1.2545851	2.187158e-01
trt4	0.3666667	0.2922613	1.2545851	2.187158e-01
trt1:plot2	0.3333333	0.2922613	1.1405319	2.625293e-01
trt2:plot2	-0.4666667	0.2922613	-1.5967447	1.201523e-01
trt3:plot2	-0.5000000	0.2922613	-1.7107978	9.679780e-02
trt4:plot2	1.3000000	0.2922613	4.4480744	9.800566e-05
trt1:plot3	-0.9000000	0.2922613	-3.0794361	4.236568e-03
trt2:plot3	1.4000000	0.2922613	4.7902340	3.652662e-05
trt3:plot3	-1.8333333	0.2922613	-6.2729254	4.937112e-07
trt4:plot3	0.1000000	0.2922613	0.3421596	7.344680e-01
trt1:plot4	-0.2000000	0.2922613	-0.6843191	4.987001e-01
trt2:plot4	0.8666667	0.2922613	2.9653829	5.673462e-03

# correct analysis: random-effects model

Use the `lme4` package for random effects. `nlme` is an alternative, but can only handle nested random effects, not crossed random effects.

```
library(lme4)
fit2 = lmer(y ~ trt + (1 | trt:plot), dat)
summary(fit2)
```

```
Linear mixed model fit by REML ['lmerMod']
Formula: y ~ trt + (1 | trt:plot)
Data: dat

REML criterion at convergence: 73.5

Scaled residuals:
    Min       1Q   Median       3Q      Max
-1.47874 -0.77009  0.08077  0.61410  1.44532
```

```
Random effects:
Groups   Name             Variance Std.Dev.
trt:plot (Intercept) 0.4402    0.6635
Residual                0.1281    0.3579
Number of obs: 48, groups:  trt:plot, 16
```

```
Fixed effects:
              Estimate Std. Error t value
(Intercept)  12.0750    0.3475   34.751
trt2           0.7417    0.4914    1.509
trt3          -0.9750    0.4914   -1.984
trt4           1.0583    0.4914    2.154
```

```
Correlation of Fixed Effects:
      (Intr) trt2    trt3
trt2 -0.707
trt3 -0.707  0.500
trt4 -0.707  0.500  0.500
```

note the same estimate for the 2 variance components as estimated above:

- 0.4402 for plot variation
- 0.1281 for subsample variation (between “cores”)

Now to test treatment effects:

```
anova(fit2) # tests fixed effects
```

```
Analysis of Variance Table
      npar Sum Sq Mean Sq F value
trt       3  2.6006  0.86687   6.7658
```

```
# then manual calculation of p-value based on 16-4 = 12 df denominator:
pf(6.7658, df1=3, df2=12, lower.tail=F) # p-value = 0.006365
```

```
[1] 0.006365647
```

and to test for variation between plots:

```
fit2.null = lm(y ~ trt, dat)
anova(fit2, fit2.null) # put complex model first!
```

```
refitting model(s) with ML (instead of REML)
```

```
Data: dat
Models:
fit2.null: y ~ trt
fit2: y ~ trt + (1 | trt:plot)
      npar      AIC      BIC logLik deviance Chisq Df Pr(>Chisq)
fit2.null    5 107.635 116.991 -48.818   97.635
fit2         6  83.795  95.022 -35.898   71.795 25.84  1 3.709e-07
```

Warning: the test above is a LRT test. Downsides:

- it's only approximate, and
- it's conservative because  $\sigma^2=0$  is at the boundary of the parameter space.

Alternative, use tools from the `lmerTest` package to get p-values:

```
library(lmerTest)
fit3 = lmer(y ~ trt + (1 | trt:plot), dat)
anova(fit3)
```

```
Type III Analysis of Variance Table with Satterthwaite's method
      Sum Sq Mean Sq NumDF DenDF F value    Pr(>F)
trt  2.6006  0.86687      3     12  6.7658 0.006366
```

```
drop1(fit3)
```

Single term deletions using Satterthwaite's method:

```
Model:
y ~ trt + (1 | trt:plot)
      Sum Sq Mean Sq NumDF DenDF F value    Pr(>F)
trt  2.6006  0.86687      3     12  6.7658 0.006366
```

```
ranova(fit3) # to test for random effects, but LRT. F-test better when appropriate
```

ANOVA-like table for random-effects: Single term deletions

Model:

```
y ~ trt + (1 | trt:plot)
      npar  logLik      AIC      LRT Df Pr(>Chisq)
<none>      6 -36.752  85.503
(1 | trt:plot)  5 -51.634 113.267 29.764  1  4.88e-08
```

## pool subsamples?

Here we get something equivalent to the correct mixed-model analysis because the design is balanced.

Note that the result of the tests (f statistic & p-value) are the same as before, but the sums of squares for `trt` and residual SS are exactly 1/3 as large as before, because there are only 1/3 as many data points used in the analysis.

```
dat_byplot = dat %>% group_by(trt:plot) %>% summarize(trt=trt[1], plot=plot[1], y=mean(y))
head(as.data.frame(dat_byplot), n=5)
```

	trt:plot	trt	plot	y
1	1:1	1	1	12.26667
2	1:2	1	2	12.60000
3	1:3	1	3	11.36667
4	1:4	1	4	12.06667
5	2:1	2	1	12.36667

```
fit4 = lm(y ~ trt, dat_byplot)
anova(fit4)
```

Analysis of Variance Table

Response: y

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
trt	3	9.8024	3.2675	6.7658	0.006366
Residuals	12	5.7953	0.4829		

```
summary(fit4)$coefficients
```

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	12.0750000	0.3474694	34.751266	2.055069e-13
trt2	0.7416667	0.4913959	1.509306	1.570928e-01
trt3	-0.9750000	0.4913959	-1.984144	7.058967e-02
trt4	1.0583333	0.4913959	2.153729	5.229395e-02

What we lack from this analysis is the ability to measure how subsampling helps. With this analysis, we can't use this experiment to guide future experiments in terms of how many subsamples are optimal.

# treatment differences and contrasts

The output of the `lmer` fit shows the SEs to compare pairs of treatments (0.4914) and the SEs to get the confidence interval for a single treatment mean (0.3475 –here for treatment 1, but it's the same for all treatment means because the design is balanced).

```
# fit2@beta    # same values as below (fixed effects), no names
# fixef(fit2) # fixed effects
summary(fit2)$coefficients
```

	Estimate	Std. Error	t value
(Intercept)	12.0750000	0.3474694	34.751264
trt2	0.7416667	0.4913959	1.509306
trt3	-0.9750000	0.4913959	-1.984144
trt4	1.0583333	0.4913959	2.153728

`lmerTest` also makes it very easy to make pairwise comparisons (but warning: *no* multiple comparison protection here)

```
ls_means(fit3)
```

Least Squares Means table:

	Estimate	Std. Error	df	t value	lower	upper	Pr(> t )
trt1	12.07500	0.34747	12	34.751	11.31793	12.83207	2.055e-13
trt2	12.81667	0.34747	12	36.886	12.05960	13.57374	1.011e-13
trt3	11.10000	0.34747	12	31.945	10.34293	11.85707	5.588e-13
trt4	13.13333	0.34747	12	37.797	12.37626	13.89040	7.562e-14

Confidence level: 95%

Degrees of freedom method: Satterthwaite

```
ls_means(fit3, which="trt", pairwise=TRUE)
```

Least Squares Means table:

	Estimate	Std. Error	df	t value	lower	upper	Pr(> t )
trt1 - trt2	-0.741667	0.491396	12	-1.5093	-1.812326	0.328993	0.157093
trt1 - trt3	0.975000	0.491396	12	1.9841	-0.095660	2.045660	0.070590
trt1 - trt4	-1.058333	0.491396	12	-2.1537	-2.128993	0.012326	0.052294
trt2 - trt3	1.716667	0.491396	12	3.4934	0.646007	2.787326	0.004435
trt2 - trt4	-0.316667	0.491396	12	-0.6444	-1.387326	0.753993	0.531427
trt3 - trt4	-2.033333	0.491396	12	-4.1379	-3.103993	-0.962674	0.001376

Confidence level: 95%

Degrees of freedom method: Satterthwaite

or we can use the `emmeans` package

```
library(emmeans)
fit3_em = emmeans(fit3, "trt")
fit3_em
```

trt	emmean	SE	df	lower.CL	upper.CL
1	12.1	0.347	12	11.3	12.8
2	12.8	0.347	12	12.1	13.6
3	11.1	0.347	12	10.3	11.9
4	13.1	0.347	12	12.4	13.9

Degrees-of-freedom method: kenward-roger  
Confidence level used: 0.95

```
pairs(fit3_em) # Tukey correction by default, unlike above
```

contrast	estimate	SE	df	t.ratio	p.value
1 - 2	-0.742	0.491	12	-1.509	0.4621
1 - 3	0.975	0.491	12	1.984	0.2468
1 - 4	-1.058	0.491	12	-2.154	0.1916
2 - 3	1.717	0.491	12	3.493	0.0200
2 - 4	-0.317	0.491	12	-0.644	0.9155
3 - 4	-2.033	0.491	12	-4.138	0.0065

Degrees-of-freedom method: kenward-roger  
P value adjustment: tukey method for comparing a family of 4 estimates

```
pwpm(fit3_em, adjust="none", means=F) # no correction: LSD. bottom: t-values
```

	1	2	3	4
1		0.1571	0.0706	0.0523
2	-0.742		0.0044	0.5314
3	0.975	1.717		0.0014
4	-1.058	-0.317	-2.033	

Row and column labels: trt  
Upper triangle: P values  
Lower triangle: Comparisons (estimate)    earlier vs. later

Next, let's test the contrast  $\mu_1 + \mu_2 + \mu_3 - 3\mu_4$ , say. The `contrast` function from `emmeans` is the easiest to use, in terms of setting up our coefficients:

```
res = contrast(fit3_em, list(trt123_vs_4 = c(1,1,1,-3)))
# summary(res)$estimate; summary(res)$t.ratio # to get more precision
res
```

```
contrast      estimate    SE df t.ratio p.value
trt123_vs_4    -3.41  1.2 12  -2.832  0.0151
```

Degrees-of-freedom method: kenward-roger

or we can use the `contest` function from `lmerTest`, which requires being careful when setting up our coefficients:

```
fixef(fit3) # to get the meaning and order of coefficients
```

```
(Intercept)      trt2      trt3      trt4
12.0750000    0.7416667  -0.9750000   1.0583333
```

```
# sum(c(0,1,1,-3) * fixef(fit3)) # contrast value
contest(fit3, c(0,1,1,-3)) # "con"trast "test", from lmerTest library
```


```
Sum Sq Mean Sq NumDF DenDF F value Pr(>F)
1 1.027315 1.027315     1    12  8.018069 0.0151302
```

Note the F-value above is  $8.020224 = (-2.832)^2 = t^2$ .

## check assumptions

A plot of the total residuals that combine plot and subsample variation is *not* very helpful:

```
ggplot(data.frame(trt=dat$trt, resid=resid(fit2), plot=dat$plot),
  aes(x=trt, y=resid, color=plot, shape=plot)) +geom_point() +theme_minimal()
```



Let's first check the assumptions on the residuals at the subsample level, i.e. the term that is not accounted for in the model (the last residual level of variation).

For this, it will be useful to get the model predictions that do or do not use the estimated random effects:




```
head(cbind(
  predict(fit2), # uses estimated random effects (re) by default
  predict(fit2, re.form=NA) )) # prediction for new plots: unknown random effects
```


```
      [,1]      [,2]
1 12.24972 12.075
2 12.24972 12.075
3 12.24972 12.075
4 12.55357 12.075
5 12.55357 12.075
6 12.55357 12.075
```

To get the residuals at the subsample level, we calculate the difference between the observations and the predictions that know about the estimated plot (random) effects.

```
subsample_residuals = dat$y - predict(fit2)
tmp = data.frame(fitted=predict(fit2), resid=subsample_residuals, trt=dat$trt)
ggplot(tmp, aes(x=fitted, y=resid, color=trt, shape=trt)) +geom_point() +theme_minimal()
```



```
ggplot(tmp, aes(sample=resid)) + geom_qq() + theme_minimal()
```




To look at residuals at the level of plots: we can extract the best linear unbiased estimates (BLUPs) of the random effects. In the residual plots below, we should have 16 points only: 1 per plot.


```
re = ranef(fit2); head(re$trt, n=4)
```

```
(Intercept)
1:1  0.174716806
1:2  0.478572120
1:3 -0.645692543
1:4 -0.007596383
```

```
tmp = data.frame(resid = re[[1]][,1],
                  trt = substr(rownames(re[[1]]), 1,1))
tmp$fitted = predict(fit2, tmp, re.form=NA)
ggplot(tmp, aes(x=fitted, y=resid)) + geom_point() + theme_minimal()
```




```
ggplot(tmp, aes(sample=resid)) + geom_qq() + theme_minimal()
```



```
library(lattice) # for dotplot function
# BLUPs and 95% prediction intervals for random effects for plots
dotplot(ranef(fit2, condVar=TRUE))[[1]]
```

## trt:plot



Conclusion: at both levels (plots and subsamples), the assumption of a normal distribution and equal variance seems adequate.